Title: NewsBot Intelligence System 2.0 — Reflective Journal
Name: Mohanad Yassin
Course: ITAI2373 —Final Project
Date: 12/01/2025

I handled this project all by myself. I did not work with any other people. Still, I managed it well. I kept all the parts organized. I made sure everything was well documented. When I ran into problems, I looked up online tutorials and other resources. That was sort of like having invisible helpers showing me the next steps.

NewsBot turned out to be tough to build. It was worth the effort though. The biggest early issue came with preparing the dataset. I had to clean up those BBC News articles. I removed any missing data. I also worked to balance the categories across the board. That process ate up a good amount of time. After that, I moved on to building TF-IDF features. I needed them to really highlight the differences between the various categories. Picking suitable classification models was another tough spot. I tried Logistic Regression. I used Naive Bayes. Random Forest came into play too. Then I evaluated each one carefully. I wanted to pick the strongest option. Beyond that, I dealt with pulling out named entities. I ran sentiment analysis right alongside it. That meant juggling outputs from spaCy and TextBlob with extra care. Finally, I had to integrate all eight modules into a single notebook. Those included preprocessing. TF-IDF was there. POS tagging fit in. Syntax parsing worked next. Sentiment analysis rounded out part of it. Classification handled the core. NER wrapped it up. Getting them to flow together without issues took some serious planning. I went through plenty of tests and fixes along the way.

Even with all those hurdles, the whole project turned out to be rewarding in the end. NewsBot handles raw articles pretty well now. It cleans them up first. Then it sorts them into groups like Politics, Sports, Technology, Business, Entertainment, and Health. It picks up on sentiment and emotions too. It pulls out main entities such as people, organizations, and locations. It even spots patterns in how things are written and what words get used. A setup like this seems useful for keeping track of news trends. It could help with digging into business details. Or it might just clarify whats going on in various fields.

I ended up managing the entire thing on my own. That meant I dealt with cleaning and setting up the dataset. I put together all the NLP parts and ran tests on them. I made some visualizations to show results. I wrote up notes in the notebook to cover everything. I also created the README file. And I put this reflection together. Taking care of it all by myself let me grasp each stage really deeply. It built up my assurance that I could handle a full NLP effort starting from nothing.

Things could improve for NewsBot down the line in a couple of ways. Switching to stronger NLP tools like BERT or RoBERTa might boost how accurate it gets with emotions. Adding more data from extra news outlets and topics would strengthen the whole system. Building an interactive dashboard would let users check results live more easily. Linking up entities to show connections could reveal even more from the information.